Find the Promos

Problem 1 · Screenplay-to-promo shortlisting

Why this problem and scope

Editors scan entire series to find a handful of moments worth testing. I built a tool that narrows the search to continuous 20–60-second cuts from a single episode. The screenplay supplies timings and visible action. Editors check the footage, select the cuts and export an edit sheet.

LangGraph workflow

Runs once per episode after ZIP parsing and scene annotation.

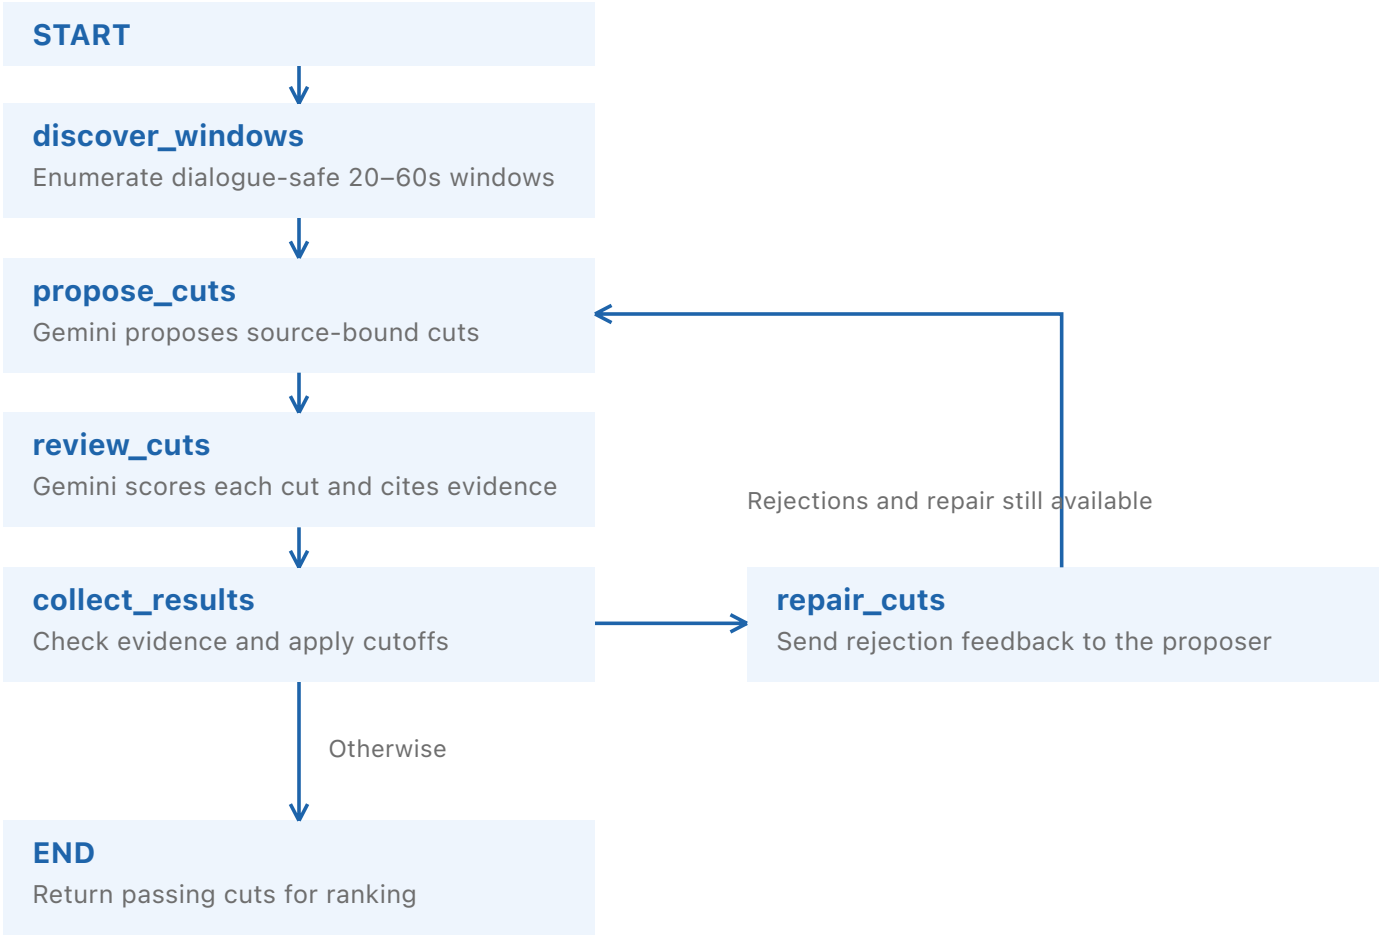

Gemini awards hook /30, clarity /25, cliffhanger /35, and pacing /10. Python applies minima of 25, 20, 28, and 7, then ranks and deduplicates across the project.

Results and quality

50	25	244
episodes completed	ranked promos	model attempts

Full-pack run

The run recorded 28 passing candidates and 72 saved rejections. Deduplication removed three overlapping alternatives. The final shortlist covers 23 episodes and exceeds my soft target of about 20 promos. Nine ranked promos tied at 93, which limits separation at the top.

What I would post first

Episode 32, 01:59.50–02:34.38, scored 93. Catherine raises a branding poker near the opening. The cut ends with its glowing tip close to Juliana’s eye and leaves the physical consequence unresolved.

Where it picked badly

Episode 7, 01:05.20–01:31.50, also scored 93. Gerald turns to leave before the cut ends, weakening the immediate cliffhanger despite its 32/35 score. In episode 16, Gemini proposed a 12-second cut, which the duration validator blocked. These checks use the screenplay. Audience performance remains unmeasured.

Scale and cost

Episodes can be queued and checkpointed independently, with concurrency limited by provider quotas. The run averaged about five model attempts per episode, with five recoverable validation failures in total. For production, measure input and output tokens and apply the configured model’s rates. Track editor acceptance and time saved.

AI tools used

I used ChatGPT/Codex for prototyping, debugging, and tests, and Gemini 3.1 Flash-Lite for scene analysis, cut proposals, and review. Episode 7 shows where the model’s judgment needs editor review.